

Dynamic Indexes

Advanced Programming and Algorithmic Design

Alberto Casagrande

Email: `acasagrande@units.it`

a.a. 2018/2019

The background of the slide features a large, faint watermark of the University of Trieste logo. The logo is circular and contains the text "UNIVERSITA' DEGLI STUDI DI TRIESTE" around the perimeter. In the center, there is an illustration of a building with a dome and a tower, with the words "E SPLENDI" below it.

Motivations

A Simple Problem for Registry Office

Let us consider the registry office

For each newborn, they record a set of data e.g., name, birthday, parents, etc.

So, the registry data-set (hopefully) changes quite often

What if they frequently perform a birthday-based search on the data-set? E.g., Find all the baby born a given (variable) day?

Some Possible Strategies

They may:

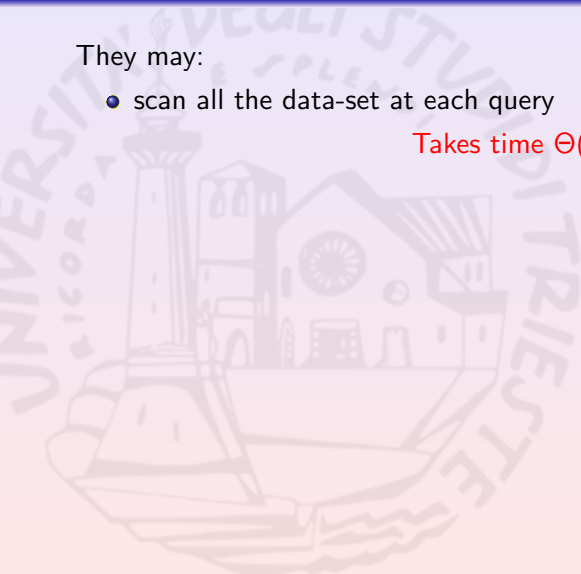


Some Possible Strategies

They may:

- scan all the data-set at each query

Takes time $\Theta(n)$



Some Possible Strategies

They may:

- scan all the data-set at each query

Takes time $\Theta(n)$

- sort the data-set by birthday at each insertion

Takes time $\Theta(n)$ per insertion (Radix Sort) + $\Theta(\log n)$ per query. What if they also want to perform searches by name?

Some Possible Strategies

They may:

- scan all the data-set at each query

Takes time $\Theta(n)$

- sort the data-set by birthday at each insertion

Takes time $\Theta(n)$ per insertion (Radix Sort) + $\Theta(\log n)$ per query. What if they also want to perform searches by name?

- sort the data-set by birthday at each query

Takes time $\Theta(n + \log n) = \Theta(n)$ per query
(Radix Sort and DS)

Some Possible Strategies

They may:

- scan all the data-set at each query

Takes time $\Theta(n)$

- sort the data-set by birthday at each insertion

Takes time $\Theta(n)$ per insertion (Radix Sort) + $\Theta(\log n)$ per query. What if they also want to perform searches by name?

- sort the data-set by birthday at each query

Takes time $\Theta(n + \log n) = \Theta(n)$ per query
(Radix Sort and DS)

The data-set often changes, thus, array is not the most suitable data structure to achieve this goal

A Dynamic Data Structure for Indexing

We need a data structure providing (efficient) support for:

- adding new data
- searching data
- removing data

A Dynamic Data Structure for Indexing

We need a data structure providing (efficient) support for:

- adding new data
- searching data
- removing data

We are aiming to build an **index** i.e., an auxiliary data structure to “efficiently” perform above operations

Binary Search Tree Notation

We have already introduced trees as ADT.

As far as searching trees concern, the following notation holds

Binary Search Tree Notation

We have already introduced trees as ADT.

As far as searching trees concern, the following notation holds

`x.left` and `x.right` are the left and right children of `x`, respectively

We have already introduced trees as ADT.

As far as searching trees concern, the following notation holds

`x.left` and `x.right` are the left and right children of `x`, respectively

`x.parent` is the parent of x

Binary Search Tree Notation

We have already introduced trees as ADT.

As far as searching trees concern, the following notation holds

`x.left` and `x.right` are the left and right children of `x`, respectively

`x.parent` is the parent of `x`

`x.key` is the value stored in `x` i.e., its key

`x.left=NIL`, then `x` misses the left child

We have already introduced trees as ADT.

As far as searching trees concern, the following notation holds

`x.left` and `x.right` are the left and right children of `x`, respectively

`x.parent` is the parent of x

x.key is the value stored in x i.e., its key

x.left=NIL, then x misses the left child

T.root is the root of the tree T and $T.root.parent = \text{NIL}$

```
def IS_RIGHT_CHILD(x):
    return x.parent != NIL and x.parent.right == x
enddef

def SIBLING(x):    # get x's sibling
    if IS_RIGHT_CHILD(x):
        return x.left
    endif
    return x.right
enddef

def UNCLE(x):      # get x's uncle
    return SIBLING(x.parent)
enddef
```

Some Useful $\Theta(1)$ Functions (Cont'd)

```

def CHILDHOOD_SIDE(x): # get x's side w.r.t.
                        # its parent
    if IS_RIGHT_CHILD(x):
        return RIGHT
    endif
    return LEFT
enddef

def REVERSE_SIDE(side): # reverse the side
    if side = LEFT:
        return RIGHT
    endif
    return LEFT
enddef

```

Some Useful $\Theta(1)$ Functions (Cont'd 2)

```
def GET_CHILD(x, side):  # get x's child on side
    if side = LEFT:
        return x.left
    endif

    return x.right
enddef

def SET_CHILD(x, side, y):  # set x's child
    if side = LEFT:
        x.left ← y
    else:
        x.right ← y
    endif
enddef
```


Some Useful $\Theta(1)$ Functions (Cont'd 3)

```
def GRANDPARENT(x):    # get x's granpa  
    return x.parent.parent  
enddef
```

Binary Search Trees

A **Binary Search Tree (BST)** is a tree s.t.:


- all the keys belong to a totally ordered set w.r.t. \preceq



Binary Search Trees

A **Binary Search Tree (BST)** is a tree s.t.:


- all the keys belong to a totally ordered set w.r.t. \preceq
- if x_l is in the left sub-tree of x , then $x_l.key \preceq x.key$




Binary Search Trees

A **Binary Search Tree (BST)** is a tree s.t.:


- all the keys belong to a totally ordered set w.r.t. \preceq
- if x_l is in the left sub-tree of x , then $x_l.key \preceq x.key$
- if x_r is in the right sub-tree of x , then $x.key \preceq x_r.key$



Binary Search Trees: an Example



Binary Search Trees: an Example




In-Order Walk

```
def INORDER_WALK_AUX(x):  
    if x ≠ NIL:  
        INORDER_WALK_AUX(x.left)  
        print x.key  
        INORDER_WALK_AUX(x.right)  
    endif  
endif  
  
def INORDER_WALK(T):  
    INORDER_WALK_AUX(T.root)  
endif
```

Searching for the Maximum/Minimum

Due to the BST property:


- the minimum key is contained by the first node on the leftmost branch that has not a left child



Searching for the Maximum/Minimum

Due to the BST property:


- the minimum key is contained by the first node on the leftmost branch that has not a left child



Searching for the Maximum/Minimum

Due to the BST property:


- the minimum key is contained by the first node on the leftmost branch that has not a left child
- the maximum key is contained by the first node on the rightmost branch that has not a right child



Searching for the Maximum/Minimum

Due to the BST property:

- the minimum key is contained by the first node on the leftmost branch that has not a left child
- the maximum key is contained by the first node on the rightmost branch that has not a right child



Searching for the Maximum/Minimum: Pseudo-Code

```
def MINIMUM_IN_TREE(x):  
    while x.left  $\neq$  NIL:  
        x  $\leftarrow$  x.left  
    endif  
    return x  
endif
```

```
def MAXIMUM_TREE(x):  
    while x.right  $\neq$  NIL:  
        x  $\leftarrow$  x.right  
    endif  
    return x  
endif
```


Searching for the Maximum/Minimum: Pseudo-Code

```
def MINIMUM_IN_TREE(x):  
    while x.left ≠ NIL:  
        x ← x.left  
    endif  
    return x  
endif
```


$O(h_T)$

```
def MAXIMUM_TREE(x):  
    while x.right ≠ NIL:  
        x ← x.right  
    endif  
    return x  
endif
```

$O(h_T)$

Successor of a Node


Due to the BST property, the successor w.r.t. \preceq of n is either



Successor of a Node

Due to the BST property, the successor w.r.t. \preceq of n is either


- the node containing the minimum in the left sub-tree of n



Successor of a Node

Due to the BST property, the successor w.r.t. \preceq of n is either


- the node containing the minimum in the left sub-tree of n or



Successor of a Node

Due to the BST property, the successor w.r.t. \preceq of n is either


- the node containing the minimum in the left sub-tree of n or
- the nearest “left-ancestor” of n , if n has not left child



Successor of a Node

Due to the BST property, the successor w.r.t. \preceq of n is either

- the node containing the minimum in the left sub-tree of n or
- the nearest “left-ancestor” of n , if n has not left child



Successor of a Node: Pseudo-Code

```
def SUCCESSOR(x):  
    if x.left  $\neq$  NIL:  
        return SEARCH_MIN_SUBTREE(x)  
    endif  
  
    y  $\leftarrow$  x.parent  
    while y  $\neq$  NIL and IS_RIGHT_CHILD(x):  
        x  $\leftarrow$  y  
        y  $\leftarrow$  x.parent  
    endwhile  
  
    return y  
enddef
```

Successor of a Node: Pseudo-Code

```
def SUCCESSOR(x):  
    if x.left ≠ NIL:  
        return SEARCH_MIN_SUBTREE(x)  
    endif  
  
    y ← x.parent  
    while y ≠ NIL and IS_RIGHT_CHILD(x):  
        x ← y  
        y ← x.parent  
    endwhile  
  
    return y  
enddef
```


$O(h_T)$

Searching for a Value in a BST

To search for v , apply a dichotomic approach from the root n :

- if n is NIL or $n.key = v$, return n
- if $n.key \leq v$, search on the right sub-tree
- if $n.key \not\leq v$, search on the left sub-tree

Searching for 3




Searching for a Value in a BST

To search for v , apply a dichotomic approach from the root n :

- if n is NIL or $n.key = v$, return n
- if $n.key \leq v$, search on the right sub-tree
- **if $n.key \not\leq v$, search on the left sub-tree**

Searching for 3




Searching for a Value in a BST

To search for v , apply a dichotomic approach from the root n :

- if n is NIL or $n.key = v$, return n
- **if $n.key \leq v$, search on the right sub-tree**
- if $n.key \not\leq v$, search on the left sub-tree

Searching for 3




Searching for a Value in a BST

To search for v , apply a dichotomic approach from the root n :

- if n is NIL or $n.key = v$, return n
- **if $n.key \leq v$, search on the right sub-tree**
- if $n.key \not\leq v$, search on the left sub-tree

Searching for 3




Searching for a Value in a BST

To search for v , apply a dichotomic approach from the root n :

- if n is NIL or $n.key = v$, return n
- if $n.key \leq v$, search on the right sub-tree
- if $n.key \not\leq v$, search on the left sub-tree

Searching for 3



Searching for a Value in a BST: Pseudo-Code

```
def SEARCH_SUBTREE(x, v):  
    if x = NIL:  
        return x  
    endif  
  
    if x.key  $\preceq$  v:  
        if v  $\preceq$  x.key:  
            return x  
        endif  
  
        return SEARCH_SUBTREE(x.right, v)  
    endif  
  
    return SEARCH_SUBTREE(x.left, v)  
enddef
```

Searching for a Value in a BST: Complexity

At each iteration, algorithm performs $\Theta(1)$ operations


The # of iterations depends on the height h_T of T and on v

The algorithm takes time $O(h_T)$

Inserting a Value in a BST

Browse a branch of T and add a new leaf having v as key


Inserting 3



Inserting a Value in a BST

Browse a branch of T and add a new leaf having v as key


Inserting 3



Inserting a Value in a BST

Browse a branch of T and add a new leaf having v as key


Inserting 3



Inserting a Value in a BST

Browse a branch of T and add a new leaf having v as key


Inserting 3



Inserting a Value in a BST

Browse a branch of T and add a new leaf having v as key


Inserting 3



Inserting a Value in a BST

Browse a branch of T and add a new leaf having v as key

Inserting 3



Inserting a Value in a BST: Pseudo-Code

```

def INSERT_BST(T,z): # z is the new node
  x ← T.root
  y ← NIL    # y is x's parent

  # search the right place for z
  while x ≠ NIL:
    y ← x
    if z.key ≤ x.key:
      x ← x.left
    else:
      x ← x.right
    endif
  endwhile

```

Inserting a Value in a BST: Pseudo-Code

```
def INSERT_BST(T, z): # z is the new node  
  x ← T.root  
  y ← NIL    # y is x's parent  
  
  # search the right place for z  
  while x ≠ NIL:  
    y ← x  
    if z.key ≤ x.key:  
      x ← x.left  
    else:  
      x ← x.right  
    endif  
  endwhile
```

 $O(h_T)$

Inserting a Value in a BST: Pseudo-Code (Cont'd)

attaching the new node

`z.parent ← y`

if `y=NIL`:

`T.root ← z`

else:

if `z.key ≤ y.key`:

`y.left ← z`

else:

`y.right ← z`

endif

endif

enddef

Inserting a Value in a BST: Pseudo-Code (Cont'd)

attaching the new node

`z.parent ← y`

if `y=NIL`:

`T.root ← z`

else:

if `z.key ≤ y.key`:

`y.left ← z`

else:

`y.right ← z`

endif

endif


enddef

$\Theta(1)$

Removing a Key from a BST

We want to remove z . Either


- z has one child at most



Removing a Key from a BST

We want to remove z . Either


- z has one child at most



Removing a Key from a BST

We want to remove z . Either


- z has one child at most



Removing a Key from a BST

We want to remove z . Either


- z has one child at most



Removing a Key from a BST

We want to remove z . Either


- z has one child at most



Removing a Key from a BST

We want to remove z . Either


- z has one child at most



Removing a Key from a BST

We want to remove z . Either


- z has one child at most or
- z has two children



Removing a Key from a BST

We want to remove z . Either


- z has one child at most or
- z has two children



Removing a Key from a BST

We want to remove z . Either


- z has one child at most or
- z has two children



Removing a Key from a BST

We want to remove z . Either


- z has one child at most or
- z has two children



Removing a Key from a BST

We want to remove z . Either


- z has one child at most or
- z has two children



Removing a Key from a BST

We want to remove z . Either


- z has one child at most or
- z has two children



Removing a Key from a BST

We want to remove z . Either


- z has one child at most or
- z has two children



Removing a Key from a BST

We want to remove z . Either


- z has one child at most or
- z has two children



Removing a Key from a BST

We want to remove z . Either

- z has one child at most or
- z has two children



Removing a Key from a BST: Pseudo-Code

```

def TRANSPLANT(T,x,y): # replace x by y
    if x.parent=NIL: # x is T's root
        T.root ← y
    else:             # x has a parent
        x_side ← CHILDHOOD_SIDE(x)
        GET_CHILD(x.parent, x_side) ← y # attach y in place of x
    endif

    if y≠NIL:         # update y's parent
        y.parent ← x.parent
    endif
enddef

```

Removing a Key from a BST: Pseudo-Code (Cont'd)

```

def REMOVE_BST(T,z): # remove z.key from T and
                    # return a removed node
    if z.left=NIL:   # if z lacks of left child
        TRANSPLANT(T,z,z.right)
        return z
    endif

    if z.right=NIL:  # if z lacks of right child
        TRANSPLANT(T,z,z.left)
        return z
    endif

    y ← MINIMUM_IN_TREE(z.right)
    z.key ← y.key

    return REMOVE_BST(T,y) # y lacks of left child
enddef

```

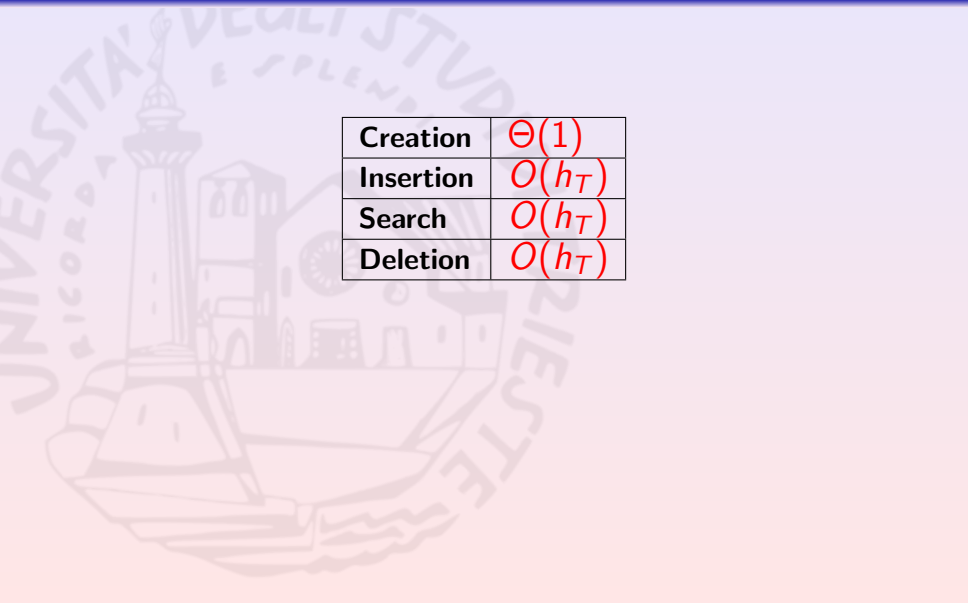
Remove a Node from a BST: Complexity

TRANSPLANT costs $\Theta(1)$, while MINIMUM_IN_TREE $O(h_T)$

So, if z has at most one child, removing it costs $\Theta(1)$

In the general case, removing z takes time $O(h_T)$

To Summarize BSTs...



Creation	$\Theta(1)$
Insertion	$O(h_T)$
Search	$O(h_T)$
Deletion	$O(h_T)$

To Summarize BSTs...

Creation	$\Theta(1)$
Insertion	$O(h_T)$
Search	$O(h_T)$
Deletion	$O(h_T)$

However, h_T may be equal to the number n of nodes e.g., keep inserting always the maximum

To Summarize BSTs...

Creation	$\Theta(1)$
Insertion	$O(n)$
Search	$O(n)$
Deletion	$O(n)$

However, h_T may be equal to the number n of nodes e.g., keep inserting always the maximum

BSTs are exceeded by single-linked lists (insertion $\Theta(1)$)

Balancing BSTs

The minimum height for a binary tree having n nodes is $\lceil \log_2 n \rceil$

We aim to balance the trees we dealing with i.e., bring their heights to $O(\log n)$

How to do it?

Balancing BSTs

The minimum height for a binary tree having n nodes is $\lceil \log_2 n \rceil$

We aim to balance the trees we dealing with i.e., bring their heights to $O(\log n)$

How to do it?

GLOBALLY: balance the whole BST after each insertion/deletion

Balancing BSTs

The minimum height for a binary tree having n nodes is $\lceil \log_2 n \rceil$

We aim to balance the trees we dealing with i.e., bring their heights to $O(\log n)$

How to do it?

GLOBALLY: balance the whole BST after each insertion/deletion
 $O(n)$

Balancing BSTs

The minimum height for a binary tree having n nodes is $\lceil \log_2 n \rceil$

We aim to balance the trees we dealing with i.e., bring their heights to $O(\log n)$

How to do it?

GLOBALLY: balance the whole BST after each insertion/deletion
 $O(n)$

LOCALLY: balance only the “unbalanced” part of the tree

Balancing BSTs

The minimum height for a binary tree having n nodes is $\lceil \log_2 n \rceil$

We aim to balance the trees we dealing with i.e., bring their heights to $O(\log n)$

How to do it?

GLOBALLY: balance the whole BST after each insertion/deletion
 $O(n)$

LOCALLY: balance only the “unbalanced” part of the tree
How to know if it is unbalanced?

Balancing BSTs

The minimum height for a binary tree having n nodes is $\lceil \log_2 n \rceil$

We aim to balance the trees we dealing with i.e., bring their heights to $O(\log n)$

How to do it?

GLOBALLY: balance the whole BST after each insertion/deletion
 $O(n)$

LOCALLY: balance only the “unbalanced” part of the tree
How to know if it is unbalanced? How to handle branches' lengths?

The background of the slide features a large, faint watermark of the University of Pisa logo. The logo is circular and contains an illustration of the Leaning Tower of Pisa. The text "UNIVERSITA' DEGLI STUDI DI PISA" is written around the top half of the circle, and "FONDATA NEL 1543" is at the bottom. The word "PISTE" is visible on the right side of the circle.

Red-Black Trees

RBTs: Definition

Are BSTs satisfying the following conditions:

- each node is either a RED or a BLACK node
- the tree's root is BLACK
- all the leaves are BLACK NIL nodes
- all the RED nodes must have BLACK children
- for each node n , all the branches from n contain the same # of black nodes

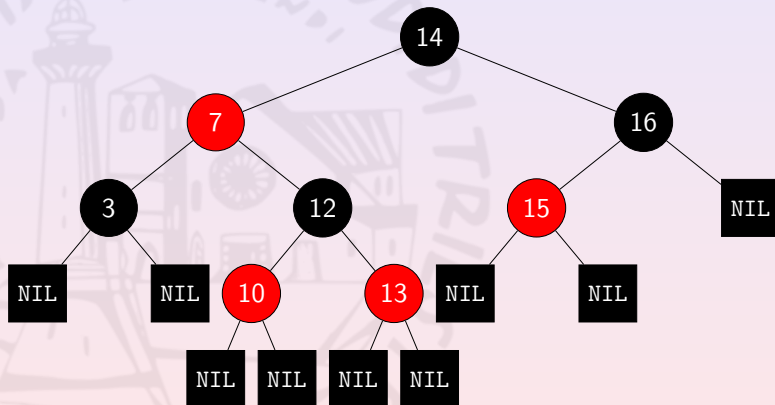
RBTs: Definition

Are BSTs satisfying the following conditions:

- each node is either a RED or a BLACK node
- the tree's root is BLACK
- all the leaves are BLACK NIL nodes
- all the RED nodes must have BLACK children
- for each node n , all the branches from n contain the same # of black nodes

$BH(x)$ will be the # of BLACK nodes $\neq x$ in any branch from x

RBTs: An Example



How “Tall” Are RB-Trees?

Theorem (Heights of a RB-Tree)

Any RBT with n internal nodes has height at most $2 \log_2(n + 1)$

Proof Sketch:

- prove that the sub-tree rooted in x has at least $2^{BH(x)} - 1$ internal nodes
- $BH(x)$ is at least half of x 's height h then

$$n \geq 2^{h/2} - 1$$

How “Tall” Are RB-Trees?

The ratio between x 's height and $BH(x)$ is topped by 2

Theorem (Heights of a RB-Tree)

Any RBT with n internal nodes has height at most $2 \log_2(n + 1)$


Proof Sketch:

- prove that the sub-tree rooted in x has at least $2^{BH(x)} - 1$ internal nodes
- $BH(x)$ is at least half of x 's height h then

$$n \geq 2^{h/2} - 1$$


Rotating a Sub-Tree

Rotations are operations on the tree structure preserving the binary searching tree property




Rotating a Sub-Tree

Rotations are operations on the tree structure preserving the binary searching tree property



Rotating a Sub-Tree

Rotations are operations on the tree structure preserving the binary searching tree property



Rotating a Sub-Tree: Pseudo-Code

```

def ROTATE(T,x,side):
    r_side ← REVERSE_SIDE(side)

    y ← GET_CHILD(x, r_side)
    TRANSPLANT(T,x,y)

    beta ← GET_CHILD(y, side)
    TRANSPLANT(T,beta,x)

    SET_CHILD(x,r_side,beta) # move beta
    if beta ≠ NIL:
        beta.parent ← x
    endif
enddef

```

Rotating a Sub-Tree: Pseudo-Code

```

def ROTATE(T, x, side):
    r_side ← REVERSE_SIDE(side)

    y ← GET_CHILD(x, r_side)
    TRANSPLANT(T, x, y)

    beta ← GET_CHILD(y, side)
    TRANSPLANT(T, beta, x)

    SET_CHILD(x, r_side, beta) # move beta
    if beta ≠ NIL:
        beta.parent ← x
    endif
enddef

```


 $\Theta(1)$

Inserting a New Node


Requires:

- inserting as in BST
- RED coloring the node
- fixing up RB-Tree properties


Inserting a New Node: Example



Inserting a New Node: Example




Inserting a New Node: Example




Inserting a New Node: Example

z's parent may be RED. How to fix it?




Inserting a New Node: Example

Case 1: z 's uncle (y) is RED...




Inserting a New Node: Example

Case 1: z 's uncle (y) is RED... RED color z 's granpa and BLACK color z 's parent and y .




Inserting a New Node: Example

Case 1: z 's uncle (y) is RED...




Inserting a New Node: Example

Still facing problems and z 's uncle is BLACK (no Case 1)




Inserting a New Node: Example

Case 2: y is BLACK and y and z are both right children.




Inserting a New Node: Example

Case 2: y is BLACK and y and z are both right children. Rotate left on z 's parent.




Inserting a New Node: Example

Case 2: y is BLACK and y and z are both right children. Rotate left on z 's parent. New z is former z 's parent




Inserting a New Node: Example

Still facing problems, but




Inserting a New Node: Example

Still facing problems, but y is still BLACK (no Case 1)




Inserting a New Node: Example

Still facing problems, but y is still BLACK (no Case 1) and z and y are on different sides w.r.t. their parents (no Case 2)




Inserting a New Node: Example

Case 3: y is BLACK and y and z are right and left children.




Inserting a New Node: Example

Case 3: y is BLACK and y and z are right and left children.
Invert z 's pa and granpa colors




Inserting a New Node: Example

Case 3: y is BLACK and y and z are right and left children.
Invert z 's pa and granpa colors




Inserting a New Node: Example

Case 3: y is BLACK and y and z are right and left children.
Invert z 's pa and granpa colors and rotate right on z 's granpa



Inserting a New Node: Example

Case 3: y is BLACK and y and z are right and left children.
Invert z 's pa and granpa colors and rotate right on z 's granpa



Inserting a New Node: Complexity

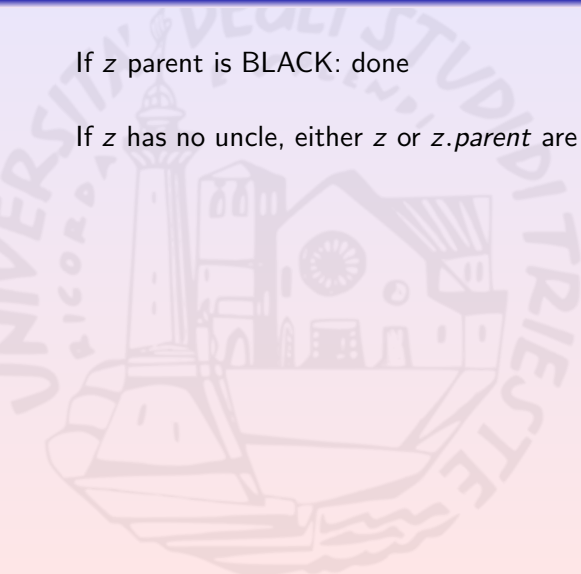
If z parent is BLACK: done



Inserting a New Node: Complexity

If z parent is BLACK: done

If z has no uncle, either z or $z.parent$ are the root: BLACK color it



Inserting a New Node: Complexity

If z parent is BLACK: done

If z has no uncle, either z or $z.parent$ are the root: BLACK color it

If z has an uncle, we can always choose between Cases 1, 2, or 3

Inserting a New Node: Complexity

If z parent is BLACK: done

If z has no uncle, either z or $z.parent$ are the root: BLACK color it

If z has an uncle, we can always choose between Cases 1, 2, or 3

Case 1 either removes the problem or pushes it towards the root

Inserting a New Node: Complexity

If z parent is BLACK: done

If z has no uncle, either z or $z.parent$ are the root: BLACK color it

If z has an uncle, we can always choose between Cases 1, 2, or 3

Case 1 either removes the problem or pushes it towards the root

Case 2 brings to Case 3

Inserting a New Node: Complexity

If z parent is BLACK: done

If z has no uncle, either z or $z.parent$ are the root: BLACK color it

If z has an uncle, we can always choose between Cases 1, 2, or 3

Case 1 either removes the problem or pushes it towards the root

Case 2 brings to Case 3

Case 3 solves the problem

Inserting a New Node: Complexity

If z parent is BLACK: done

If z has no uncle, either z or $z.parent$ are the root: BLACK color it

If z has an uncle, we can always choose between Cases 1, 2, or 3

Case 1 either removes the problem or pushes it towards the root

Case 2 brings to Case 3

Case 3 solves the problem

In the worst case, the algorithm keeps repeating Case 1 steps along the insertion branch and the complexity is $O(\log n)$

Inserting a New Node: Code

```
def INSERT_RBTREE(T, z):  
    INSERT_BST(T, z)  
    z.color ← RED  
  
    FIX_INSERT_RBTREE(T, z)  
enddef  
  
def FIX_INSERT_RBT_CASE1(T, z):  
    UNCLE(z).color ← BLACK  
    z.parent.color ← BLACK  
    GRANDPARENT(z).color ← RED  
  
    return GRANDPARENT(z)  
endif
```

Inserting a New Node: Code (Cont'd)

```
def FIX_INSERT_RBT_CASE2(T, z):  
    p ← z.parent  
  
    z_side ← CHILDHOOD_SIDE(z)  
    ROTATE(T, p, REVERSE_SIDE(z_side))  
  
    return p  
endif
```

Inserting a New Node: Code (Cont'd 2)

```
def FIX_INSERT_RBT_CASE3(T, z):  
    g ← GRANDPARENT(z)  
  
    z.parent.color ← BLACK  
    g.color ← RED  
  
    ROTATE(T, g, CHILDHOOD_SIDE(z))  
endif
```

Inserting a New Node: Code (Cont'd 3)

```

def FIX_INSERT_RBTREE(T, z):
    while (z.parent ≠ NIL and
           (GRANDPARENT(z) ≠ NIL or z.parent.color = RED)):
        if UNCLE(z).color = RED:
            z ← FIX_INSERT_RBT_CASE1(T, z)
        else:
            if (CHILDHOOD_SIDE(z) ≠
                CHILDHOOD_SIDE(z.parent)):
                z ← FIX_INSERT_RBT_CASE2(T, z)
            endif
            FIX_INSERT_RBT_CASE3(T, z)
        endif
    endwhile

    T.root.color ← BLACK
enddef

```

Removing a Key

Removing a key as in BST removes also a node y which is replaced by its former child x



Removing a Key

Removing a key as in BST removes also a node y which is replaced by its former child x

If y was RED, the RB-Tree properties are preserved

Removing a Key

Removing a key as in BST removes also a node y which is replaced by its former child x

If y was RED, the RB-Tree properties are preserved

If y was BLACK, the branches through x lost 1 BLACK node

In particular, $BH(x) = BH(w) - 1$ where w is x 's sibling

Removing a Key

Removing a key as in BST removes also a node y which is replaced by its former child x

If y was RED, the RB-Tree properties are preserved


If y was BLACK, the branches through x lost 1 BLACK node

In particular, $BH(x) = BH(w) - 1$ where w is x 's sibling

The fixing procedure iteratively balances BH on the sub-tree rooted on x 's parent

Removing a Key: Case 0

x is RED




Removing a Key: Case 0

x is RED


- BLACK color x

$BH(x)$ is increased by 1 and the tree has been fixed



Removing a Key: Case 1

x 's sibling is RED




Removing a Key: Case 1

x 's sibling is RED


- invert colors in x 's parent and sibling
- rotate x 's parent on x 's side

$BH(x)$ does not change



Removing a Key: Case 2

x 's sibling and nephews are BLACK




Removing a Key: Case 2

x 's sibling and nephews are BLACK


- RED color x 's sibling

$BH(x)$ does not change, while the BLACK height of both x 's parent and sibling are decreased by 1



Removing a Key: Case 3

Among x 's sibling and nephews, only the nephew on x 's side is RED




Removing a Key: Case 3

Among x 's sibling and nephews, only the nephew on x 's side is RED


- rotate x 's sibling on the opposite side w.r.t. x
- invert colors in both old and new siblings of x

The BLACK height of both x and x 's parent does not change



Removing a Key: Case 4

The x 's nephew on the opposite side w.r.t. x is RED




Removing a Key: Case 4

The x 's nephew on the opposite side w.r.t. x is RED


- switch colors of x 's parent and sibling
- BLACK color the x 's nephew on the opposite side w.r.t. x
- rotate x 's parent on x 's side

The BLACK height of both x and x 's parent does not change



Removing a Key: Some Considerations

All the case transformation procedures takes time $\Theta(1)$



Removing a Key: Some Considerations

All the case transformation procedures takes time $\Theta(1)$

Both Case 0 and Case 4 transformation procedures fix the RB-Tree

Removing a Key: Some Considerations

All the case transformation procedures takes time $\Theta(1)$

Both Case 0 and Case 4 transformation procedures fix the RB-Tree

Case 1 cannot occur twice in-row

Removing a Key: Some Considerations

All the case transformation procedures takes time $\Theta(1)$

Both Case 0 and Case 4 transformation procedures fix the RB-Tree

Case 1 cannot occur twice in-row

Case 2 pushes the problem one step towards the root

Removing a Key: Some Considerations

All the case transformation procedures takes time $\Theta(1)$

Both Case 0 and Case 4 transformation procedures fix the RB-Tree

Case 1 cannot occur twice in-row

Case 2 pushes the problem one step towards the root

If Case 2 occurs after Case 1, Case 0 occurs next

Removing a Key: Some Considerations

All the case transformation procedures takes time $\Theta(1)$

Both Case 0 and Case 4 transformation procedures fix the RB-Tree

Case 1 cannot occur twice in-row

Case 2 pushes the problem one step towards the root

If Case 2 occurs after Case 1, Case 0 occurs next

Case 3 transformation procedure brings to Case 4

Removing a Key: Some Considerations

In the worst case scenario, Case 2 is repeated until the problem is pushed up to the root ($O(\log n)$ times)

If this is the case, we have decreased the BLACK height of the tree and the problem is no more a problem

Removing a Key: Code

```
def REMOVE_RBTREE(T, z):  
    y ← REMOVE_BST(T, z)  
  
    if y.color = BLACK: # fix from its replacement  
        if y.left = NIL  
            x ← y.right  
        else:  
            x ← y.left  
        endif  
        FIX_REMOVE_RBTREE(T, x)  
    endif  
  
    return y  
enddef
```

Removing a Key: Code (Cont'd 2)

```
def FIX_REMOVE_RBT_CASE1(T, x):  
    SIBLING(x).color ← BLACK  
    x.parent.color ← RED  
  
    ROTATE(T, x, CHILDHOOD_SIDE(x))  
endif  
  
def FIX_REMOVE_RBT_CASE2(T, x):  
    SIBLING(x).color ← RED  
    return x.parent  
endif
```

Removing a Key: Code (Cont'd 3)

```

def FIX_REMOVE_RBT_CASE3(T,x):
    x_side ← CHILDHOOD_SIDE(x)
    r_side ← REVERSE_SIDE(x_side)

    w ← GET_CHILD(w, r_side)
    GET_CHILD(w, x_side).color ← BLACK
    w.color ← RED

    ROTATE(T,w, r_side)
endif

```

Removing a Key: Code (Cont'd 4)

```

def FIX_REMOVE_RBT_CASE4(T, x):
    x_side ← CHILDHOOD_SIDE(x)
    r_side ← REVERSE_SIDE(x_side)

    w ← GET_CHILD(w, r_side)
    GET_CHILD(w, r_side).color ← BLACK
    w.color ← x.parent.color
    x.parent.color ← BLACK

    ROTATE(T, x.parent, x_side)
endif

```

Removing a Key: Code (Cont'd 5)

```

def FIX_REMOVE_RBT(T, x):
    while x ≠ T.root and x.color ≠ RED:
        w ← SIBLING(x)
        if w = RED:
            x ← FIX_REMOVE_RBT_CASE1(T, x)
        else:
            x_side ← CHILDHOOD_SIDE(x)
            r_side ← REVERSE_SIDE(x_side)

```


Removing a Key: Code (Cont'd 6)

```

if GET_CHILD(w, r_size) = RED:
    FIX_REMOVE_RBT_CASE4(T, x)

```

```

    return

```

```

else :

```

```

    if GET_CHILD(w, x_side) = RED:
        FIX_REMOVE_RBT_CASE3(T, x)

```

```

    else :

```

```

        FIX_REMOVE_RBT_CASE2(T, x)

```

```

    endif

```

```

endif

```

```

endwhile

```

```

enddef

```